Small Quantum Classifiers Under Matched Classical and Hardware Tests:

Pre-Registered Benchmarks, Fourier Surrogates and IBM Heron Runs

Vivek Reddy
Independent researcher
vvardhan00@gmail.com

November 2026

Abstract

We compare small quantum neural networks (QNNs) with classical models under a fixed protocol: classical baselines get the same inputs and are tuned by cross-validation, hardware runs are compared with noise-free simulation on identical inputs, and hardware deviations are tested against the exact shot-noise expectation. In a pre-registered evaluation on untouched data, a 4-qubit QNN on sentence-encoder features is within about two points of tuned classical models on the same features and has no higher point estimate than any of them (sentiment 0.766 vs. 0.769; CVE severity 0.708 vs. 0.720, a tie within a pre-set ±3-point margin); a 1.5-point gap seen during design did not replicate. An ablation crossing encoding with re-uploading attributes most of the design effect to the encoding (19–24 points). The circuit's output lies exactly in a small Fourier basis; a classical model trained on that basis outperforms the QNN on CVE (0.727 vs. 0.708) and is inconclusive on sentiment (0.756 vs. 0.766). On IBM's `ibm_marrakesh` (200 inputs per task; three CVE runs, one sentiment run), no run's accuracy differs significantly from noise-free, but the mean output shifts towards one class in every run (-0.002 to -0.006). Single observed executions on two further Heron devices show the same sign with larger shifts (to -0.015); the device's noise model matched `ibm_marrakesh` but understated `ibm_fez`. In simulation, a trained two-copy QNN separates pure from mixed states with 16–32 copies per state from 2 to 6 qubits, against 128–4,096 for the classical-shadow baseline we test, illustrating a known multi-copy separation; in a preliminary 2-qubit hardware run on one shared 60-state set its copy ratio is 2.5–2.9 (job-conditional intervals above 1), with point estimates below the same-state noise-free ratio (6.3) but intervals too broad to establish a reduction. All code, results and job identifiers are archived.

1 Introduction

Claims of quantum advantage in machine learning are hard to evaluate when results are reported without matched classical baselines [1]. We ask a narrow question: for small, trainable quantum classifiers that run on current hardware, on which tasks do they differ from well-tuned classical models, and how much does real-device noise change their outputs?

We report two projects. The first trains a 4-qubit QNN on classical data, including two text-classification tasks, and runs the trained circuits on IBM Quantum hardware. The second classifies quantum states directly, a setting where theory predicts separations for learners with quantum memory [2–4]; its purity result is in exact simulation up to 6 qubits, with a 2-qubit check on hardware. Our contributions are:

- an evaluation protocol that keeps quantum and classical models on identical inputs, tunes the classical baselines, and compares hardware with simulation on identical inputs (Section 2);
- evidence from a pre-registered evaluation on untouched data, with seed variation and paired intervals, that a well-designed QNN is close to but not better than tuned classical models on the same features for two text tasks, and an ablation of the circuit design (Section 3);
- a repeated hardware study on fixed physical qubits, with paired tests against noise-free simulation, showing no significant accuracy change but a small output bias of consistent sign and varying size, replicated in sign on two further devices (Section 4);
- a small-scale simulation of the copy-complexity gap between a learned two-copy QNN and a classical-shadow baseline for purity testing, with per-seed results, measured gradient variance, and a 2-qubit hardware run in which the gap narrows but persists (Section 5).

2 Evaluation protocol

1. **Same inputs.** Classical baselines (logistic regression and an RBF SVM; an MLP where noted) receive exactly the features encoded into the circuit.
2. **Tuned baselines and a classical reference.** Same-feature baselines are tuned by 5-fold cross-validation on the training data only. We also report logistic regression on the full feature set and all training data.
3. **Same inputs on hardware.** Device runs use a fixed random subset of the test set; they are compared only with the noise-free simulator on that subset.
4. **Shot noise as the null hypothesis.** Prediction changes on hardware are compared with the exact distribution expected from finite sampling alone, and repeated in separate jobs.
5. **Pre-registered confirmation.** Because the design was chosen with test-set feedback (below), the final text-task comparison is repeated once on untouched data, with the system, the analysis and the decision rules committed publicly before that data was downloaded.

Disclosure. The circuit design and feature-reduction method were chosen in exploratory runs that reported test-set accuracy; no separate validation set was held out for those choices. The test-set numbers in Table 1 are therefore exploratory; the pre-registered evaluation in Table 2 is the confirmatory result for the text tasks. The ablation and the quantum-data experiments have no such untouched replication. Rule 3 was introduced after our first hardware chart placed 40-input device accuracy beside 200-input simulator accuracy, which made the device appear better than the simulator on one dataset.

3 Quantum classifiers on classical data

Model. Inputs are reduced to one feature per qubit and min–max scaled to $[0, \pi]$ using the training data. The circuit repeats, twice, an $R_Y(x_i)$ encoding layer followed by one trainable `RealAmplitudes` layer (data re-uploading [5, 6]); the class is the measured value of qubit 0 (class 1 if $P(1) > 0.5$). Weights are initialised uniformly in $[-\pi, \pi]$ from the seed and trained with L-BFGS-B (at most 60 iterations, one start) on the cross-entropy of the exact simulated output probabilities, with one-hot labels. We used Qiskit 1.4 [7] and qiskit-machine-learning 0.8.

Text data and partitions. Texts are embedded with all-MiniLM-L6-v2 [8] (384 dimensions) and reduced to four features with partial least squares (PLS), fitted on the full training split; the scalers are also fitted on training data only. *Sentiment*: Rotten Tomatoes sentence polarity [9], the standard 8,530/1,066 train/test split, balanced; the 1,066-sentence validation split was never used until the confirmatory evaluation. *CVE severity*: 12,116 NVD descriptions with CVSS v3.1 scores, published 1 January–29 April 2024, labelled high/critical if the base score is at least 7.0, and split by publication date: training up to 9 April 2024 (9,692; 49.9% high/critical), test after it (2,424; 42.2%). The label shift explains the 0.578 majority-class baseline. For each seed, the QNN and the same-feature baselines are trained on the same balanced random subset of 600 training texts; the full-embedding reference uses all training texts.

Model (test accuracy, mean ± std over 3 seeds)	Sentiment (1,066)	CVE severity (2,424)
Majority class	0.500	0.578
QNN, 4 qubits	0.757 ± 0.007	0.699 ± 0.015
Logistic regression, same 4 features, defaults	0.770 ± 0.000	0.688 ± 0.004
Logistic regression, same 4 features, CV-tuned	0.772 ± 0.002	0.689 ± 0.004
RBF SVM, same 4 features, CV-tuned	0.770 ± 0.004	0.691 ± 0.008
Logistic regression, full embedding, all data (defaults / CV-tuned)	0.774 / 0.774	0.715 / 0.702
QNN − tuned logistic regression, 95% CI over seeds and items	[−0.029, −0.003]	[−0.013, +0.035]

Table 1: Text tasks, exploratory test sets (population std over seeds). The interval for the difference is a crossed two-level percentile bootstrap: each replicate resamples the three seeds with replacement and, independently, the test items with replacement, applying the same item resample to every selected seed (10,000 replicates); per-seed paired intervals and exact McNemar tests are in the results files. Tuning grids: $C \in \{0.01, \ldots, 100\}$; for the SVM also $\gamma \in \{\text{scale}, 0.1, 1, 10\}$.

On the exploratory test sets the QNN is about 1.5 points below tuned logistic regression on sentiment (interval excludes zero) and nominally ahead on CVE severity (interval includes zero). Cross-validated tuning of the full-embedding model lowered its CVE test accuracy (0.715 to 0.702), consistent with the label shift between the time-ordered splits; we report the tuned value as the reference.

Confirmatory evaluation. We then froze the system and scored it once on data no design decision had touched: the Rotten Tomatoes validation split, and all 13,432 CVEs with CVSS v3.1 scores published 1 May–28 August 2024 (47.9% high/critical). The protocol, script and decision rules were committed before the data was downloaded (Appendix A). The pre-registered primary contrast is the QNN minus CV-tuned logistic regression on the same features, averaged over seeds, with a 95% interval from a crossed two-level percentile bootstrap: each replicate resamples the three seeds with replacement and, independently, the test items with replacement, applying the same item resample to every selected seed (10,000 replicates). The pre-registered rules were: on sentiment, the expected gap is confirmed if that interval lies entirely below zero; on CVE severity, a tie is confirmed if it lies entirely within $[-0.03, +0.03]$ (an equivalence margin in the spirit of two one-sided tests). Contrasts against the SVM and the full-embedding reference were declared secondary. Table 2 gives the outcome. On sentiment the expected gap did not replicate: the QNN is within 0.3 points of tuned logistic regression, and the interval includes zero. On CVE severity the pre-registered tie is confirmed, but the QNN is now nominally below rather than ahead, and in a secondary contrast it is nominally 1.8 points below the tuned SVM (interval excludes zero; not adjusted for multiple comparisons). Taken together, on untouched data the QNN is within about

two points of tuned classical models on the same four features and has no higher point estimate than any of them; its only interval excluding zero is the secondary, unadjusted contrast with the tuned SVM on CVE. The seed component of these intervals rests on only three training seeds, so it is coarse. The per-seed results (in the results files) are the more direct evidence about training variation, and they show that the order of the QNN and tuned logistic regression changes between seeds on both tasks. Unrounded differences are in the results files.

Confirmatory set (accuracy, mean ± std over 3 seeds)	Sentiment (1,066)	CVE severity (13,432)
Majority class	0.500	0.521
QNN, 4 qubits	0.766 ± 0.003	0.708 ± 0.009
Logistic regression, same 4 features, CV-tuned	0.769 ± 0.006	0.720 ± 0.004
RBF SVM, same 4 features, CV-tuned	0.767 ± 0.002	0.726 ± 0.005
Logistic regression, full embedding, CV-tuned	0.777	0.725
QNN − tuned logistic regression [95% CI]	−0.003 [−0.017, +0.011]	−0.011 [−0.024, +0.003]
QNN − tuned SVM [95% CI]	−0.001 [−0.013, +0.012]	−0.018 [−0.031, −0.003]
Pre-registered verdict	expected gap not confirmed	tie within ±3 points confirmed

Table 2: Pre-registered confirmatory evaluation on untouched data (sample std over seeds; crossed two-level bootstrap intervals over seeds and items).

Dequantization check. Because each input angle is encoded twice by R_Y, whose generator has eigenvalues $\pm\frac{1}{2}$, every output of the circuit, for any weights, is a trigonometric polynomial $f(x) = \sum_\omega a_\omega \cos(\omega \cdot x) + b_\omega \sin(\omega \cdot x)$ with integer frequencies $\omega \in \{-2, \ldots, 2\}^4$ [6]: 312 frequency vectors up to sign, so 624 features plus a constant. We verified this numerically: a least-squares fit in this basis reproduces the trained seed-0 QNN on 2,000 random inputs to a maximum error of 3×10^{-15}, while a basis truncated to $\{-1, 0, 1\}^4$ misses by 0.3–0.4. We then trained purely classical surrogates on that basis [10], logistic regression with C tuned by 5-fold cross-validation, on the same features and training texts as the QNN, under a protocol committed before running. On the untouched CVE set the surrogate scores 0.727 against the QNN's 0.708 (QNN minus surrogate -0.018, $[-0.030, -0.007]$), so the QNN is classically matched and in fact exceeded. On untouched sentiment it scores 0.756 against 0.766 ($+0.010$, $[-0.008, +0.029]$), which the pre-registered rule classes as inconclusive; plain logistic regression on the four features (0.769) is ahead of both there. Random subsets of 128 frequencies perform like the full basis. Cross-validation chose the smallest C in the pre-set grid; a post-hoc rerun with a wider grid changes no result by more than 0.002. Two results should be kept apart. The representation result is exact: the circuit's output lies in this 624-feature Fourier span, so the circuit's output function is a linear combination of these classical features, whatever the trained weights. The predictive result is mixed: the trained surrogate outperforms the QNN on CVE and is inconclusive on sentiment. With 624 features and 600 training texts, cross-validation chose strong regularisation (5-fold CV accuracy at the chosen C: 0.72–0.77 per seed); the features are unscaled cosines and sines in $[-1, 1]$. A remaining sentiment edge over this particular surrogate would reflect how the two models are trained and regularised, not a quantum resource, since the function class is the same.

Which design choice mattered? An earlier design, using the `ZZFeatureMap` encoding [11], parity readout of all qubits and training on 1,024-shot sampled probabilities with COBYLA, scored about 0.69 on breast cancer. Table 3 changes one factor at a time from the final design, over three seeds (the seed sets the 70/30 split and the initial weights), and crosses encoding with re-uploading in a 2×2 design. Replacing angle encoding with `ZZFeatureMap` lowers accuracy by 24

points with re-uploading (0.910 to 0.667) and by 19 points without (0.887 to 0.700), so in both settings the encoding dominates (the encoding effect is positive in all six seed-by-setting cells, 12–30 points). Three seeds cannot estimate the interaction: it is +0.057 with a 95% t-interval over seeds of $[-0.23, +0.35]$; re-uploading itself helps angle encoding (+2.3 points) and slightly hurts `ZZFeatureMap` (−3.3). The other single changes, including 2 or 6 qubits instead of 4, move accuracy by at most 1.4 points. Breast cancer was part of the exploratory benchmark, so these numbers are descriptive.

Variant (breast cancer, 300 training / 171 test, PCA features)	Accuracy	vs. final
Final: angle encoding, re-uploaded twice, qubit-0 readout, exact, L-BFGS-B, 4 qubits	0.910 ± 0.011	–
`ZZFeatureMap` encoding, re-uploaded, rest as final	0.667 ± 0.052	−0.244
No re-uploading (angle encoding)	0.887 ± 0.015	−0.023
`ZZFeatureMap` encoding, no re-uploading	0.700 ± 0.032	−0.211
Parity readout	0.924 ± 0.005	+0.014
COBYLA instead of L-BFGS-B	0.903 ± 0.003	−0.008
Sampled (1,024-shot) training, COBYLA	0.899 ± 0.012	−0.012
Original design (all of the above changes)	0.690 ± 0.051	−0.220
2 qubits / 6 qubits	0.916 / 0.899	+0.006 / −0.012

Table 3: Ablation, mean ± std over 3 seeds (population std); the first four rows form a 2×2 design of encoding and re-uploading; all variants are evaluated with exact probabilities. This is a single dataset and should be read as indicative.

Figure 1 gives the same comparison on three standard datasets and on the two text tasks, using smaller, separate splits (300 training and at most 200 test items, PCA for the tabular data). The QNN is within 0–10 points of the best classical model on each and above logistic regression only on the nonlinear two-moons task, where the RBF SVM is best. The text-task numbers in Table 1 supersede the text-task bars in this figure.

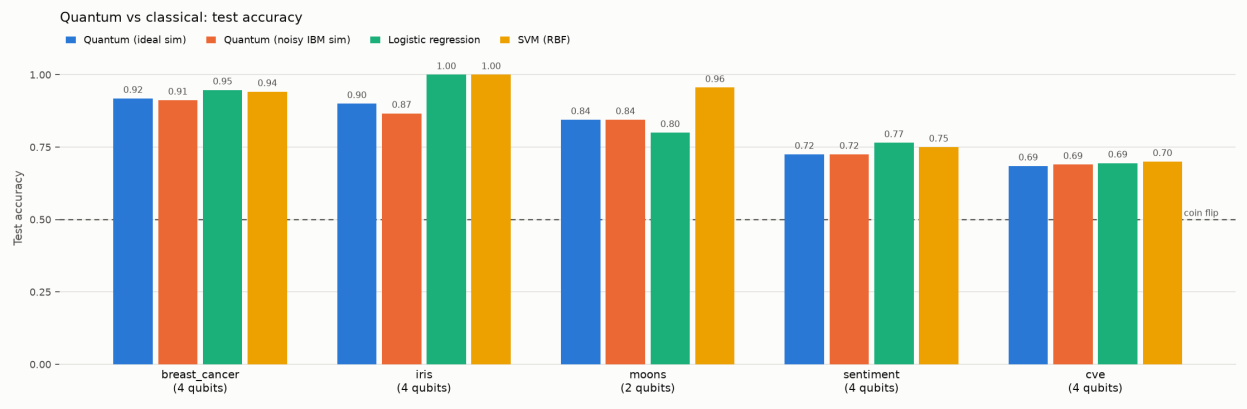


Figure 1: Benchmark on smaller, separate splits: test accuracy of the QNN (noise-free and with the `FakeManilaV2` noise model) and classical models on identical reduced features, one seed.

4 Real-hardware behaviour

Setup. We ran the trained sentiment and CVE models (seed 0, Table 1) on IBM's 156-qubit Heron device `ibm_marrakesh` through the free Open Plan, with `SamplerV2`, 1,024 shots and no error mitigation or dynamical decoupling. Each circuit was transpiled once with Qiskit's preset

pass manager (optimisation level 1, `seed_transpiler`=0), so every run used the same physical qubits (0–3); the compiled circuit has depth 37 with six CZ gates. At execution time the readout errors of these qubits were 0.006–0.024 and the CZ errors 0.0022–0.0027 (full snapshot and job identifiers in Appendix A). For each task we drew a fixed, stratified random sample of 200 test inputs and ran it as separate jobs: three for CVE severity (two at 16:16 and 16:18 UTC, and one about three hours later at 19:22 UTC) and one for sentiment. The device's reported calibration of these qubits was the same for all three CVE runs, so the later run separates them in time but not across a recalibration.

Statistics. For an input with noise-free probability p and N shots, the probability that sampling alone changes the prediction is the exact binomial tail

$$q(p) = \begin{cases} P(K \le N/2), & p > 1/2, \\ P(K > N/2), & p \le 1/2, \end{cases} \qquad K \sim \mathrm{Bin}(N, p). \tag{1}$$

Under the null hypothesis that the device adds nothing beyond shot noise, every run draws fresh shots, so outcomes are independent across inputs and runs even though runs reuse the same inputs; the total number of changed predictions then follows a Poisson-binomial distribution, which we compute exactly by convolution. A fixed device bias alone would not make runs dependent; dependence across runs would require a shared run-level effect such as drift or environmental fluctuation, which the pooled test does not model, so we also give the test for each run alone. Accuracy on the device is compared with noise-free accuracy on the same inputs by an exact McNemar test per run. Bias is the mean signed difference between device and noise-free $P(\text{class 1})$, given per run with a 95% bootstrap interval over inputs, and pooled with a two-level bootstrap that resamples runs and then inputs. The same threshold ($P > 0.5$) is used in simulation and on the device.

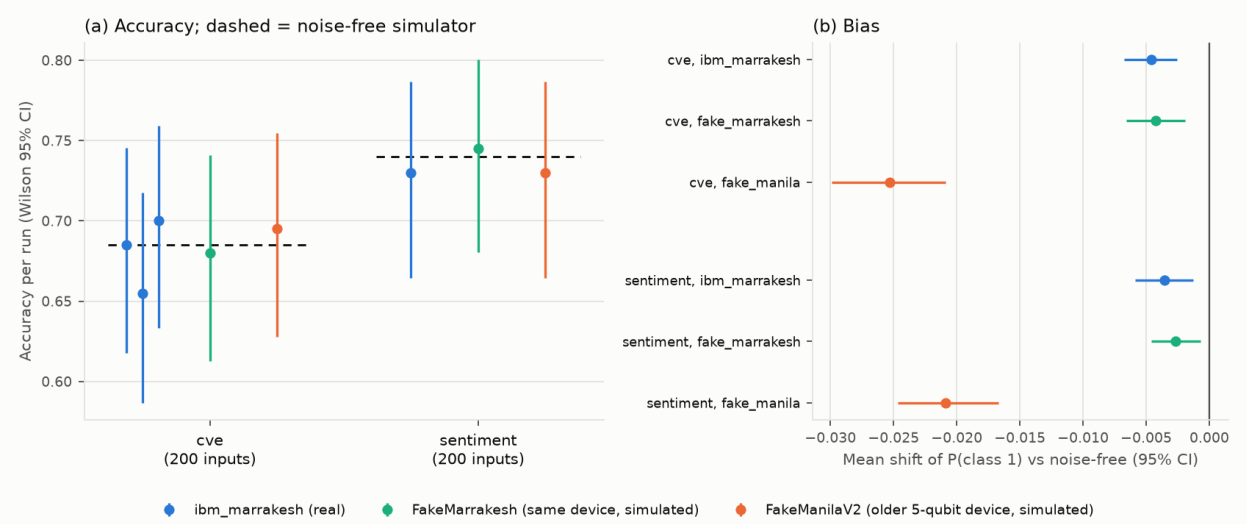


Figure 2: Hardware study, 200 inputs per task. (a) Accuracy per run with Wilson 95% intervals; dashed lines are noise-free accuracy on the same inputs. (b) Mean signed shift of $P(\text{class 1})$ relative to noise-free simulation, pooled over runs, with 95% bootstrap intervals.

Results. Figure 2 and Table 4 summarise the study. No run's accuracy differs significantly from noise-free accuracy on the same inputs: the three CVE runs score 0.685, 0.655 and 0.700 against

a noise-free 0.685 (exact McNemar $p = 1$, 0.070 and 0.45; discordant pairs 2/2, 1/7 and 5/2), and sentiment scores 0.730 against 0.740 ($p = 0.63$). With 200 inputs this rules out only large accuracy changes. The numbers of changed predictions are at the edge of what shot noise explains: 19 against 13.3 expected over three CVE runs ($p = 0.055$; per run 4, 8 and 7 against 4.4, $p = 0.68$, 0.054 and 0.13), and 4 against 3.1 for sentiment ($p = 0.37$). We read this as compatible with shot noise, not as evidence of its absence. Repeated runs agree on 96% of predictions. The mean output probability shifted towards class 0 in every run. For CVE the per-run shift was -0.0063 $[-0.0090, -0.0037]$ and -0.0053 $[-0.0081, -0.0025]$ in the first session and -0.0022 $[-0.0046, +0.0002]$ three hours later, whose interval includes zero; pooled over the three runs it is -0.0046, with a two-level interval of $[-0.0076, -0.0016]$. For sentiment (one run) it is -0.0035 $[-0.0061, -0.0010]$. The sign was the same in all four runs but the size varied between sessions, even under the same reported calibration, so we report a bias of these executions rather than a stable property of the device; the calibration snapshot records gate and readout error rates and coherence times, not drifts within a calibration cycle. Such a bias can change predictions for inputs near the decision boundary, so label agreement alone understates the effect of noise.

Task	Backend	Accuracy per run	Changed / expected	p	Bias [95% CI]
CVE (ideal 0.685)	`ibm_marrakesh`	0.685, 0.655, 0.700	19 / 13.3	0.055	-0.0046 $[-0.0076, -0.0016]$
	FakeMarrakesh	0.680	3 / 4.4	0.86	-0.0043 $[-0.0066, -0.0019]$
	FakeManilaV2	0.695	12 / 4.4	< 0.001	-0.0253 $[-0.0299, -0.0209]$
Sentiment (ideal 0.740)	`ibm_marrakesh`	0.730	4 / 3.1	0.37	-0.0035 $[-0.0059, -0.0013]$
	FakeMarrakesh	0.745	3 / 3.1	0.63	-0.0027 $[-0.0046, -0.0007]$
	FakeManilaV2	0.730	8 / 3.1	0.003	-0.0209 $[-0.0246, -0.0167]$

Table 4: Hardware study, 200 inputs per task, 1,024 shots. "Changed / expected" counts predictions that differ from the noise-free prediction, against the exact shot-noise expectation; $p = P(\text{total} \geq \text{observed})$ under shot noise alone. Bias intervals are 95% bootstraps over inputs (over runs and inputs for the three CVE runs); per-run values are in the text.

Noise models. `FakeMarrakesh`, Qiskit's noise model built from a calibration snapshot of the same device, agrees with the real device in accuracy behaviour and in the size of its bias: their difference lies within the intervals, while the bias itself is non-zero. `FakeManilaV2`, a model of an older 5-qubit device, overstates the bias about fivefold and changes significantly more predictions than shot noise allows. This compares two specific noise models on one circuit; it suggests matching the noise model to the device, rather than a general statement about noise models, and the next paragraph shows that even a same-device model can miss.

Other devices. The Open Plan gives access to two further 156-qubit Heron r2 devices, `ibm_kingston` and `ibm_fez`. With the remaining quantum time we ran the CVE model once on each, on a 100-input stratified sample that is a subset of the 200 inputs above, compiled in the same way (both used physical qubits 0–3 with the same depth-37, six-CZ circuit). Table 5 compares all three devices on these 100 inputs. No device's accuracy differs significantly from noise-free (exact McNemar $p \geq 0.13$ for every run), and the changed predictions on the two new devices are within shot noise. Every execution shifts the output towards class 0, but by different amounts: on the same inputs the `ibm_fez` execution shifts it by 0.010 more than `ibm_marrakesh` (paired 95% CI $[0.006, 0.015]$) and by 0.007 more than `ibm_kingston` ($[0.002, 0.011]$). `FakeFez` understates the shift on `ibm_fez` by 0.006 ($[0.002, 0.010]$), while `FakeMarrakesh` matches `ibm_marrakesh` (difference $+0.0005$, $[-0.003, +0.004]$); no noise model of `ibm_kingston` is distributed with Qiskit.

These intervals are over inputs, conditional on the particular executions; with one run on each new device they describe these executions, not persistent device properties. The sign of the shift was the same in every execution, its size differed between them, and a same-device noise model matched `ibm_marrakesh` but understated the `ibm_fez` execution. One run per new device cannot separate device from session effects.

Backend (100 inputs; noise-free 0.650)	Accuracy	Changed / expected	p	Mean shift [95% CI]
`ibm_marrakesh`, 3 runs	0.650, 0.600, 0.660	4, 7, 5 / 3.1	–	−0.0050 (per run −0.007, −0.007, −0.002)
`ibm_kingston`, 1 run	0.630	2 / 3.1	0.85	−0.0087 [−0.0121, −0.0055]
`ibm_fez`, 1 run	0.640	5 / 3.1	0.17	−0.0152 [−0.0208, −0.0095]
`FakeMarrakesh`	0.660	1 / 3.1	0.97	−0.0045 [−0.0078, −0.0012]
`FakeFez`	0.640	3 / 3.1	0.63	−0.0092 [−0.0135, −0.0049]

Table 5: Cross-device replication of the CVE classifier on 100 identical inputs (a subset of Table 4's inputs); `ibm_marrakesh` rows are its three runs restricted to these inputs. $p = P(\text{changed} \geq \text{observed})$ under shot noise alone.

Earlier, smaller run. A first run on 40 inputs per dataset (30 for iris) on the previous day reproduced the noise-free predictions on four of five datasets; CVE changed three predictions, and three repeats on the same 40 inputs gave 5 changes against 3.45 expected (exact $p = 0.24$; 8 against 4.6, $p = 0.063$, including the first run). These smaller samples motivated the study above.

5 Quantum classifiers on quantum data

The second project gives models direct access to quantum states; classical learners (logistic regression, RBF SVM, MLP) receive classical-shadow estimates [12] from random-Pauli snapshots, or full state vectors where noted. Each experiment uses 200 training and 400 test states per size; entanglement and phase results are over five seeds and purity results over three, each seed redrawing states, shot noise and initialisations. For the purity task, the classical learner and its hyperparameters are selected on a 50-state validation split of the training states and scored once on test. For the entanglement and phase tasks the best classical learner is selected on the test set, which favours the classical side. Throughout, "c copies" means c preparations of the state for every method: the two-copy QNN uses $c/2$ shots of a circuit holding two copies, and the shadow method uses $\lfloor c/9 \rfloor$ shots in each of 9 measurement bases.

Entanglement (2–4 qubits). The task separates product states from entangled states (Meyer–Wallach $Q \geq 0.2$). A two-copy QNN, which receives two copies of each state, is perfect in exact simulation and reaches 0.991 ± 0.006, 0.992 ± 0.004 and 0.996 ± 0.005 at 2, 3 and 4 qubits with 256 copies. A one-copy QNN stays at chance (0.47–0.52 per seed): with one copy and a fixed measurement its output probability is linear in ρ, and this label is not a linear function of the state; this does not imply that every single-copy method fails. Classical learners on full state vectors reach only 0.715 ± 0.018, 0.569 ± 0.024 and 0.524 ± 0.025, but learners on classical-shadow purity features are perfect with exact estimates and reach 0.951 ± 0.009, 0.975 ± 0.005 and 0.988 ± 0.006 with 256 snapshots. The shadow baseline therefore closes the gap to within 0.04 at 2 qubits and 0.01 at 4 qubits: this task does not separate the two approaches.

Phases. For ground states of the transverse-field Ising chain (4–8 qubits) and a $\mathbb{Z}_2$ lattice gauge model (8 link qubits), we use a QCNN with hardware-efficient units: convolution layers of $R_Y R_Z$,

CX, $R_Y R_Z$ on neighbouring qubits (8 weights per layer, shared across the ring) and coherent pooling that halves the active qubits down to one readout qubit (4 weights per layer), following the layer structure of [13]. (The 3-parameter units of the Qiskit tutorial output exactly 0.5 on these spin-flip-symmetric states.) The QCNN, a generic QNN and classical models all score 0.96–1.00 overall. These tasks are easy: logistic regression on a *single* exact feature (mean nearest-neighbour ZZ, or mean X) reaches 0.994 ± 0.005 at 4 qubits and 0.995 ± 0.006 at 8, and the mean electric field alone reaches 0.998 ± 0.002 on the gauge model. All errors fall near the critical point (about 20% of test states). There, the exact-simulation QCNN reaches 0.963 ± 0.037 (4 qubits, $|h-1| < 0.2$) and 0.974 ± 0.033 (8 qubits), against 0.901 ± 0.033 and 0.965 ± 0.014 for classical learners on exact Z/X expectations; with spreads of this size, only the 4-qubit difference is suggestive. The phase tasks do not discriminate between models.

Purity. Haar-random pure states are distinguished from random rank-2 mixed states, a global property. Single-copy strategies need a number of copies that grows exponentially with system size, while a two-copy SWAP-type test needs a constant number [2, 3]. The two-copy QNN reaches 1.000 in exact simulation at every size; the one-copy QNN stays at chance (0.44–0.57 per seed). Table 6 reports the copies of each test state needed for 90% test accuracy, on the grid $2, 4, \ldots, 4096$: 16–32 for the QNN's seed-averaged curve (16–64 for individual seeds) against 128–4,096 for the shadow baseline, and the per-seed ranges never overlap. The shadow baseline's learner (logistic regression, RBF SVM or MLP, each over a hyperparameter grid) is chosen on validation data; choosing the best learner on the test set instead gives the same thresholds except at 5 qubits (1,024 rather than 2,048).

Qubits	2	3	4	5	6
Two-copy QNN, mean curve	32	16	32	32	32
Two-copy QNN, per seed	64, 16, 16	16, 16, 16	32, 16, 16	16, 32, 16	64, 32, 16
Classical shadows (validated learner), mean curve	128	256	512	2,048	4,096
Classical shadows, per seed	128 (all)	256 (all)	512, 1024, 512	2048, 2048, 1024	4096 (all)

Table 6: Copies of each test state needed to reach 90% test accuracy, exact simulation, read off the power-of-two grid (the hardware comparison in Section 5 instead interpolates on the hardware test states). "Mean curve": smallest grid value where accuracy averaged over three seeds is at least 0.9; "per seed": the same for each seed. The one-copy QNN never reaches 90%.

Trainability. The two-copy model was trained at 2 qubits and its shared block transferred to larger sizes with 30 fine-tuning steps. Training with a separate block per copy pair from a random start failed at 5 qubits in one of three seeds (0.85 ± 0.26). To test whether this reflects a barren plateau [14, 15], we measured the variance of the gradient of the class-separation signal $D = \overline{P}_{\text{mixed}} - \overline{P}_{\text{pure}}$ at random initialisations (100 draws $\times$ 3 seeds, exact gradients, 2–7 qubits). With per-pair blocks the variance decays from 9.5×10^{-4} to 1.0×10^{-6}, a factor of 3.9 per qubit (fitted $\log_2$ slope -1.97, 95% CI over seeds $[-2.06, -1.87]$; per seed -1.95, -1.94, -2.01; close to 4^{-n}): exponential-like decay over the measured range, consistent with the barren plateaus expected for this global cost, though six sizes cannot establish the asymptotic rate. With one block shared across pairs the decay slows to a factor of 1.7 per qubit (2.6×10^{-3} to 1.6×10^{-4}; slope -0.80, 95% CI over seeds $[-1.26, -0.47]$). No initialisation in either model had a gradient norm below 10^{-3} in this range. The decay is a plausible cause of the failed seed, and weight sharing coincided with successful training, but we did not test the causal link directly.

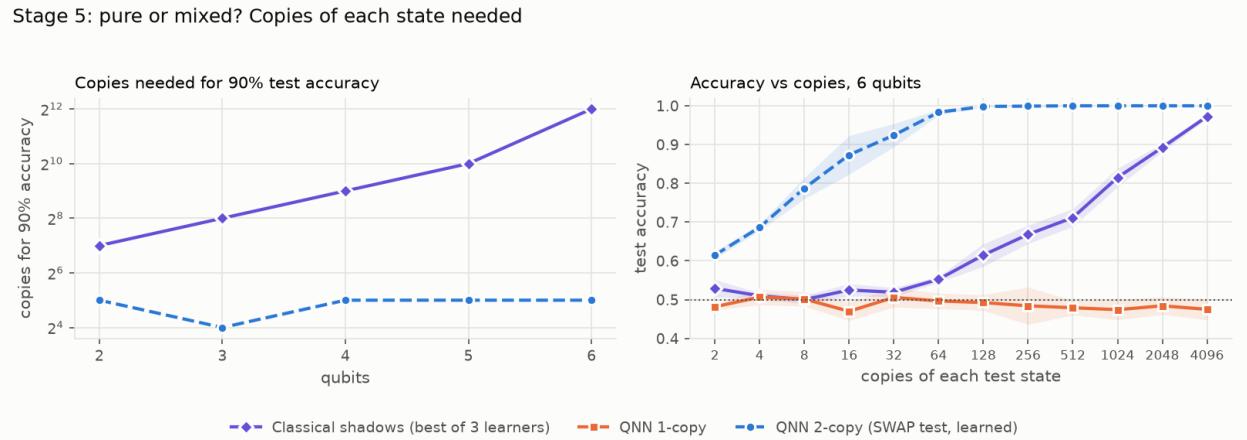


Figure 3: Pure vs. mixed states, exact simulation. Left: copies needed for 90% accuracy (mean curves). Right: accuracy vs. copies at 6 qubits (mean and spread over 3 seeds).

Purity on hardware (preliminary). We ran the 2-qubit purity test on `ibm_marrakesh` as two separate jobs on the same 60 test and 20 calibration states, with 1,024 copies per state and fixed physical qubits. Training (200), test (60) and calibration (20) states are drawn independently and balanced; no calibration state is within trace distance 0.3 of any training or test state. Mixed states are prepared as 3-qubit purifications with the ancilla unmeasured, so that pure and mixed preparations both compile to 7 CZ gates and circuit depth does not reveal the label. After routing, the two-copy circuit has about 19 CZ gates against 7 for a shadow circuit. Models are trained in simulation; on the device each method's decision rule is then refitted on the 20 calibration states ("recalibrated"): a two-parameter logistic regression on the QNN output, and the same on the shadow method's global-purity estimate, so both get the same flexibility. Recalibration changes accuracy by at most about 0.03 at any copy count, and resampling which states are used for calibration widens the intervals but does not change the ordering. Smaller copy counts are subsampled without replacement from the device shots (30 subsamples per point).

At 1,024 copies every model is at 0.99–1.00. Figure 4 shows accuracy against copies. To estimate the copies needed for 90% with honest uncertainty, we treat each job as an independent execution unit and bootstrap it 1,000 times. Each replicate resamples the test and calibration states, resamples each state's raw shots, subsamples c copies, refits the two-parameter calibration (with the same c for calibration states) and re-interpolates the 90% threshold on a log-copies scale. The QNN needs 48 $[24, 91]$ and 55 $[27, 108]$ copies in the two jobs and the shadow method 128 $[64, 256]$ in both (medians and 95% percentile intervals). That gives ratios of 2.9 $[1.3, 7.1]$ and 2.5 $[1.1, 6.3]$, above 1 in both jobs ($P(\text{ratio} < 1) = 0.005$ and 0.017). The same procedure in noise-free simulation, on the same 80 states and design, gives 16 and 108 copies (this interpolated, same-state estimate differs from the grid-based, three-seed thresholds of Table 6, 32 and 128 at 2 qubits, which use other test states; the hardware comparison uses the same-state estimate), a ratio of 6.3 $[2.5, 14.9]$. The device ratios are smaller by point estimate, but the 95% interval of their log-ratio to the noise-free ratio includes zero in both jobs ($[-2.8, 0.8]$ and $[-3.1, 0.6]$), so a reduction is not established. Calibrating with the full 1,024 copies of each calibration state instead gives device ratios of 3.1 and 3.0 with similar intervals. An earlier analysis that bootstrapped over overlapping shot subsamples gave much narrower intervals (e.g. $[2.0, 4.0]$); it understated the uncertainty, and the coarse power-of-two grid had suggested a simulated ratio of 4. With two jobs on one set of 60 states this is an illustration, not an estimate of hardware scaling. Both methods' purity estimates

correlate with the true $\mathrm{Tr}\,\rho^2$ ($r = 0.986$ for the QNN, 0.975–0.981 for shadows) and both read low by 0.06–0.10.

We compared the device with `FakeMarrakesh` by replaying the identical compiled circuits (same physical qubits, gates and shots, downloaded from the device jobs). For these circuits the noise model understates the QNN's output shift: +0.038 and +0.041 on the device against +0.022 and +0.019 in the noise model, and the noise model gives somewhat larger point-estimate ratios. For the shadow circuits device and model agree within the run-to-run spread. The two-copy circuit's greater depth is a plausible reason, but we did not vary depth in a controlled way, so this finding is limited to these compiled circuits and runs. For the shallow text classifier of Section 4 the same noise model matched the device.

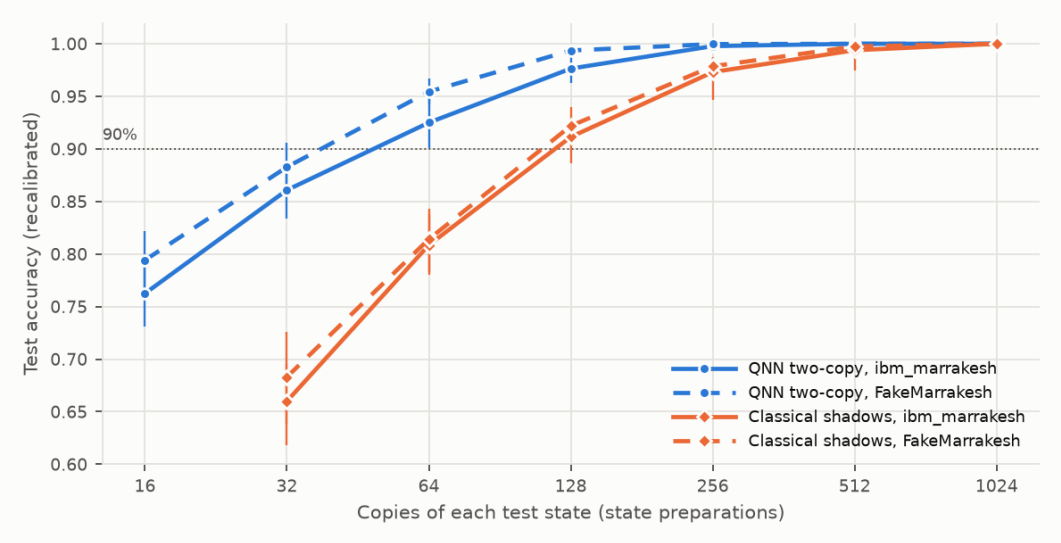


Figure 4: Pure vs. mixed, 2 qubits, on `ibm_marrakesh` (solid) and `FakeMarrakesh` replaying the identical circuits (dashed): recalibrated test accuracy against copies of each test state, subsampled from 1,024-copy jobs. Points are the mean of two jobs; bars span both jobs' 95% bootstrap intervals over the 60 test states, which do not capture variation between independent jobs. Shadows: logistic regression on the global-purity estimate. At 16 copies the shadow method has one shot per basis and was not evaluated.

Entanglement and phases on hardware. The 2-qubit entanglement and 4-qubit Ising classifiers ran as two jobs each on 40 test and 16 calibration states with 512 copies (Table 7). The two-copy QNN loses 20 points on the raw device output and recovers to 0.988 ± 0.018 after recalibration; its output shift (+0.028, +0.032) is again about twice that of `FakeMarrakesh` replaying the identical circuits (+0.016, +0.015). For the shallower QCNN device and noise model agree (-0.007 and -0.009 in both). With 80 pooled predictions a Wilson 95% interval around 0.95 spans about $[0.88, 0.98]$, so no ranking within a task is significant.

6 Related work

Benchmarking. Bowles et al. [1] benchmark QML models at scale and find that classical baselines often match or beat them; our protocol follows the same spirit at small scale, adding hardware.

Task	Model	Noise-free	Device	Recalibrated
Entanglement, 2 qubits	Two-copy QNN	1.000	0.800 ± 0.035	0.988 ± 0.018
	Best classical (snapshots)	1.000	0.938 ± 0.018	0.975 ± 0.000
Ising phases, 4 qubits	QCNN	0.950	0.925 ± 0.035	0.925 ± 0.035
	Best classical (snapshots)	0.975	0.962 ± 0.018	0.962 ± 0.018

Table 7: Quantum-data classifiers on `ibm_marrakesh`: test accuracy on 40 states, mean $\pm$ std over two runs, 512 copies. "Best classical" is the best of three learners in each column, which favours the classical side.

Encoding and expressivity. Data re-uploading [5] and the Fourier view of encoding [6] explain why the encoding dominated our ablation. **Dequantization.** Quantum-inspired classical algorithms [16] and random-Fourier-feature approximations of variational QML models [10] suggest that circuits like ours may be classically approximable; our dequantization check (Section 3) shows that our classifier's function class is exactly such a Fourier model; a classical surrogate on it outperforms the QNN on one task and is inconclusive on the other. **Learning from quantum data.** Exponential separations between learners with and without quantum memory are proven in [3, 4] and demonstrated on hardware in [2]; our purity experiment is a small-scale illustration of that separation against one classical-shadow method [12], in simulation and at 2 qubits on hardware. **Trainability.** Barren plateaus [14] and their dependence on the cost function [15] predict the exponential gradient decay we measured for per-pair blocks. **Hardware characterisation.** Capability benchmarks [17] and error mitigation [18] address device noise systematically; we used no mitigation beyond refitting decision rules on labelled device states, and characterise only output bias and accuracy for a few small circuits.

7 Discussion and limitations

In these experiments the quantum model used fewer copies than the classical baseline only in the purity task, where the tested two-copy QNN could hold two copies of a state at once, which matches theory [2, 3]; we illustrate that result rather than extend it, and only against the classical-shadow baseline we tested. On hardware the observed advantage was smaller, with broad uncertainty; the deeper two-copy circuit is a plausible contributor, and a same-device noise model understated its output shift. On classical data the circuit classifies features chosen by a classical encoder, tuned classical models on the same features do as well or slightly better, and the circuit's output lies in a Fourier basis that a classical learner can use directly, where it does at least as well on CVE.

Limitations: (i) all simulations use at most 8 qubits and are exact, so nothing here is classically hard; (ii) design choices were explored with test-set feedback (Section 2), which the confirmatory evaluation addresses only for the two text tasks; (iii) the ablation uses one dataset, itself part of the exploratory benchmark; (iv) the hardware study uses one circuit family on three devices of one processor family, with 200 inputs and few runs on one device and a single 100-input CVE run on each of the others, and the bias may drift with calibration; (v) prediction changes are tested under an independence assumption across inputs and runs; (vi) CVE severity from descriptions alone approximates CVSS scoring; (vii) training-seed variation in the text tasks is estimated from three seeds; (viii) the purity test ran on hardware only at 2 qubits and in two runs (preparation circuits for 3 or more qubits were too deep for useful fidelity, and no budget remained for a third run); (ix) the quantum-data hardware runs use 40–60 test states, too few to rank models within a task.

Reproducibility. Code, configurations, result files and job identifiers are archived on Zenodo as AGI v1.0.4 [19] (git commit `16357c3`) and Entangle v1.0.1 [20] (git commit `21ab264`), with the DOIs given in the references; every number in this paper comes from these versions. Appendix A lists software versions, commands and jobs. During this work we found that `qiskit-ibm-runtime` 0.44.1 declares Qiskit 1.x support but fails to transpile for real devices on it; we reported it upstream (Qiskit/qiskit-ibm-runtime#3563).

Acknowledgements. Code, experiments and drafting were carried out with assistance from an AI coding assistant (Claude, Anthropic); an independent critical review was obtained from another AI model (ChatGPT, OpenAI). The author reviewed all results and takes responsibility for the content. We thank IBM Quantum for Open Plan access.

References

[1] Joseph Bowles, Shahnawaz Ahmed, and Maria Schuld. Better than classical? The subtle art of benchmarking quantum machine learning models, 2024.

[2] Hsin-Yuan Huang, Michael Broughton, Jordan Cotler, Sitan Chen, Jerry Li, Masoud Mohseni, Hartmut Neven, Ryan Babbush, Richard Kueng, John Preskill, and Jarrod R. McClean. Quantum advantage in learning from experiments. *Science*, 376(6598):1182–1186, 2022. doi: 10.1126/science.abn7293.

[3] Sitan Chen, Jordan Cotler, Hsin-Yuan Huang, and Jerry Li. Exponential separations between learning with and without quantum memory. In *2021 IEEE 62nd Annual Symposium on Foundations of Computer Science (FOCS)*, pages 574–585, 2022. doi: 10.1109/FOCS52979.2021.00063.

[4] Dorit Aharonov, Jordan Cotler, and Xiao-Liang Qi. Quantum algorithmic measurement. *Nature Communications*, 13:887, 2022. doi: 10.1038/s41467-021-27922-0.

[5] Adrián Pérez-Salinas, Alba Cervera-Lierta, Elies Gil-Fuster, and José I. Latorre. Data re-uploading for a universal quantum classifier. *Quantum*, 4:226, 2020. doi: 10.22331/q-2020-02-06-226.

[6] Maria Schuld, Ryan Sweke, and Johannes Jakob Meyer. The effect of data encoding on the expressive power of variational quantum machine learning models. *Physical Review A*, 103:032430, 2021. doi: 10.1103/PhysRevA.103.032430.

[7] Ali Javadi-Abhari, Matthew Treinish, Kevin Krsulich, Christopher J. Wood, Jake Lishman, Julien Gacon, Simon Martiel, Paul D. Nation, Lev S. Bishop, Andrew W. Cross, Blake R. Johnson, and Jay M. Gambetta. Quantum computing with Qiskit, 2024.

[8] Nils Reimers and Iryna Gurevych. Sentence-BERT: Sentence embeddings using Siamese BERT-networks. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, pages 3980–3990, 2019. doi: 10.18653/v1/D19-1410.

[9] Bo Pang and Lillian Lee. Seeing stars: Exploiting class relationships for sentiment categorization with respect to rating scales. In *Proceedings of the 43rd Annual Meeting of the Association for Computational Linguistics*, pages 115–124, 2005. doi: 10.3115/1219840.1219855.

[10] Jonas Landman, Slimane Thabet, Constantin Dalyac, Hela Mhiri, and Elham Kashefi. Classically approximating variational quantum machine learning with random Fourier features, 2022.

[11] Vojtěch Havlíček, Antonio D. Córcoles, Kristan Temme, Aram W. Harrow, Abhinav Kandala, Jerry M. Chow, and Jay M. Gambetta. Supervised learning with quantum-enhanced feature spaces. *Nature*, 567:209–212, 2019. doi: 10.1038/s41586-019-0980-2.

[12] Hsin-Yuan Huang, Richard Kueng, and John Preskill. Predicting many properties of a quantum system from very few measurements. *Nature Physics*, 16:1050–1057, 2020. doi: 10.1038/s41567-020-0932-7.

[13] Iris Cong, Soonwon Choi, and Mikhail D. Lukin. Quantum convolutional neural networks. *Nature Physics*, 15(12):1273–1278, 2019. doi: 10.1038/s41567-019-0648-8.

[14] Jarrod R. McClean, Sergio Boixo, Vadim N. Smelyanskiy, Ryan Babbush, and Hartmut Neven. Barren plateaus in quantum neural network training landscapes. *Nature Communications*, 9:4812, 2018. doi: 10.1038/s41467-018-07090-4.

[15] M. Cerezo, Akira Sone, Tyler Volkoff, Lukasz Cincio, and Patrick J. Coles. Cost function dependent barren plateaus in shallow parametrized quantum circuits. *Nature Communications*, 12:1791, 2021. doi: 10.1038/s41467-021-21728-w.

[16] Ewin Tang. A quantum-inspired classical algorithm for recommendation systems. In *Proceedings of the 51st Annual ACM SIGACT Symposium on Theory of Computing*, pages 217–228, 2019. doi: 10.1145/3313276.3316310.

[17] Timothy Proctor, Kenneth Rudinger, Kevin Young, Erik Nielsen, and Robin Blume-Kohout. Measuring the capabilities of quantum computers. *Nature Physics*, 18:75–79, 2022. doi: 10.1038/s41567-021-01409-7.

[18] Kristan Temme, Sergey Bravyi, and Jay M. Gambetta. Error mitigation for short-depth quantum circuits. *Physical Review Letters*, 119:180509, 2017. doi: 10.1103/PhysRevLett.119.180509.

[19] Vivek Reddy. H3XY/AGI: quantum text classifier with fair classical baselines and IBM hardware study. Zenodo, version 1.0.4, git commit 16357c3, doi:10.5281/zenodo.24150592, 2026. URL https://doi.org/10.5281/zenodo.24150592.

[20] Vivek Reddy. H3XY/Entangle: quantum machine learning on quantum-native data, with classical-shadow baselines and IBM hardware runs. Zenodo, version 1.0.1, git commit 21ab264, doi:10.5281/zenodo.24150608, 2026. URL https://doi.org/10.5281/zenodo.24150608.

A Reproducibility details

Software. Python 3.12.16; qiskit 1.4.6; qiskit-ibm-runtime 0.43.1 (0.44.x fails on Qiskit 1.x); qiskit-machine-learning 0.8.4; qiskit-aer 0.18.2; numpy 2.4.6; scikit-learn 1.9.1; sentence-transformers 6.2.0. Training uses Qiskit's exact reference `Sampler`.

Code versions. AGI v1.0.4, commit `16357c3e38bc`, https://doi.org/10.5281/zenodo.24150592; Entangle v1.0.1, commit `21ab2642d2fb`, https://doi.org/10.5281/zenodo.24150608. Result-file paths below are relative to these versions.

Commands (AGI repository): `python train.py -task sentiment|cve`; `python -m benchmark.uncertainty` (Table 1); `python -m benchmark.ablation` (Table 3); `python -m benchmark.dequantize` (dequantization check; protocol `results/dequantization/PROTOCOL.md`, committed as `c07c087` before the first run; `-posthoc` for the wider grid); `python -m benchmark.confirmatory` (Table 2; protocol `results/confirmatory/PROTOCOL.md`, committed as `dda5f41` on 9 November 2026 at 19:37 UTC, before the confirmatory data was downloaded); `python -m benchmark.hardware_study` and `benchmark.hardware_study_report` (Table 4; `-fake-only` reproduces the noise-model rows without IBM access); `python -m ibmq.job_log`. Result files are in `results/`.

IBM jobs (Open Plan; times UTC). Cross-device runs, 12 November 2026 at 01:37–01:39 UTC (the evening of 11 November in US Central time): `ibm_kingston` `dbqhgr8k09pc73at9aig`, `ibm_fez` `dbqhh8gk09pc73at9bc0` (29 quantum seconds each; `results/hardware_study/cross_device.md`). All others on `ibm_marrakesh`. 200-input study, 9 November 2026: CVE `dbov06tt7fqc7388gab0`, `dbov3ttt7fqc7388gfrg`, `dbp1pvgou4as73cst2n0`; sentiment `dbov4itt7fqc7388ggog`. Earlier 40-input runs, 8–9 November 2026: benchmark `dbolcblt7fqc73881hv0`, `dbold1tt7fqc73881iqg`, `dboldboou4as73cs9lcg`, `dbole58qtr0c738d1ccg`, `dbolhc0qtr0c738d1ihg`; CVE repeats `dbomf5gou4as73csbdmg`, `dbomfas4148s739t8q1g`, `dbomfg5t7fqc73883ctg`. The full list with timestamps and quantum seconds is in `results/ibm_jobs.json`.

Quantum-data jobs (Entangle repository, `ibm_marrakesh`, 9 November 2026): purity `dbp0rd0ou4as73csr3t0`, `dbp0s1lt7fqc7388jap0` (pilot `dbp0r10qtr0c738dj0fg`); entanglement and Ising, run 1 `dbp11agqtr0c738djefg`, `dbp11eoou4as73csrgp0`, `dbp11mc4148s739tpbpg`, `dbp11rlt7fqc7388jns0`; run 2 `dbp124gqtr0c738djg4g`, `dbp128k4148s739tpcv0`, `dbp12g5t7fqc7388jp20`, `dbp12lc4148s739tpdm0`. Summary with Wilson intervals: `results/review_hardware.md`; seed tables `results/review_seeds_stage*.md`; validated learner selection `results/review_selection.md`; gradient variance and slope intervals `results/review_gradients.md`; hardware purity recalibration and copy accounting `results/review_purity_hw.md`, shot-level bootstrap `results/review_purity_boot.md`; per-qubit calibration of the device and of `FakeMarrakesh` `results/review_calibration.md`. Raw device counts and the compiled circuits of every job are in `results/ibm_raw/`; recomputing from them reproduces every device number, and the `FakeMarrakesh` comparisons replay those circuits. Qubits used: median T_1 205–236 μs, T_2 136–152 μs, readout error 0.007–0.018, CZ error 0.0021.

Calibration (physical qubits 0–3 at the start of the 200-input study): readout error 0.006–0.024; CZ error 0.0022–0.0027; full values per qubit, including T_1 and T_2, in `results/hardware_study/*.json`.